

How Much Future Do You Need?

Lookahead, Causality and Phone Error Rate in Streaming Accent Conversion

Dipesh Tripathi

University of South Dakota

dipesh.tripathi@coyotes.usd.edu

Abstract

Streaming foreign accent conversion (FAC) systems choose a lookahead budget and report a single end-to-end latency, but no published work sweeps that budget or separates the delay a modelling choice imposes from the delay the hardware imposes. We train a causal phone translator at 16 lookaheads from 0 to 640 ms with every other factor held fixed, and report four results. First, quality improves smoothly with lookahead at -0.045 phone error rate (PER) per doubling ($R^2 = 0.98$) with *no locatable knee*: a piecewise fit is preferred under every axis treatment we tried, yet the breakpoint moves between 40 and 180 ms and no bootstrap interval is narrower than three octaves. Instead of a knee we report a *saturation point* at ≈ 240 ms, beyond which each additional 20 ms frame buys less than the measured 2σ noise floor. A 40 ms budget, the value used by the current state of the art, forgoes $\approx 40\%$ of the achievable PER reduction. Second, a pre-registered hypothesis that vowels and approximants degrade faster than obstruents is *not supported*: all seven manner classes improve by 52–72% relative, and the pre-registered grouping explains less variance than the spread within its own groups. It still fails when each error is reweighted by its measured acoustic cost, obtained by forced-aligning per-phone spectral distortion, so the refutation is not an artefact of phone error rate charging every error equally. Third, accent conversion is not merely speech recognition with a different label: on the same dense grid, conversion buys -0.045 PER per doubling of lookahead against -0.029 for accent-faithful transcription at matched capacity, a $1.55\times$ difference, and the gap between the two widens $5.5\times$ from $L=0$ to 640 ms. Fourth, per-chunk compute is flat in lookahead (-0.3% to $+1.9\%$ across a $17\times$ change in algorithmic latency, replicated on two architectures), which makes the two latency terms independently actionable. We release the harness, and we document three measurement failures that produced confident but wrong numbers before being caught.

Index Terms: accent conversion, streaming, latency, lookahead, reproducibility

1 Introduction

Streaming FAC exists and works. PHONOS [1] runs at ≤ 40 ms lookahead and reports an 81% reduction in non-native accent confidence. What does not exist is any account of *why* 40 ms, what it costs in quality, or what any of it does on hardware that is not a datacentre GPU.

Three gaps follow, and this paper addresses the first two directly.

Nobody sweeps the budget at resolution. PHONOS states 40 ms and publishes no ablation showing 40 ms is necessary or sufficient; TVTSyn fixes 4 future tokens (≈ 80 ms) without varying them. DarkStream does ablate, over three points $\{0, 140, 280\}$ ms, and finds the third buys < 1 pp, but three points scored on anonymisation accuracy cannot locate where returns stop, which is the question a deployment budget actually asks. “What would 160 ms buy you” still has no published answer.

Algorithmic and computational latency are reported fused. A figure such as “ < 241 ms on a single GPU” combines a modelling choice with unstated inference and buffering costs. A reader cannot tell whether a faster chip would help. We decompose throughout as

$$t_{e2e} = \underbrace{c + L}_{t_{\text{alg}}} + t_{\text{cmp}} + t_{\text{buf}}, \quad (1)$$

where c is chunk size, L is lookahead, t_{cmp} is per-chunk compute and t_{buf} is I/O and jitter buffering. t_{alg} is inherent: it is the delay on an infinitely fast machine, because the model refuses to emit frame n until it has seen frame $n+L$. t_{cmp} is contingent, quantise, prune, buy a faster chip.

Degradation is reported in aggregate. Existing results are utterance-level MOS, WER or accentedness. Which *sounds* break first is unreported.

Our contributions:

1. A latency–quality characterisation of streaming FAC across 16 lookahead budgets, reporting an exchange rate and a floor-relative saturation point rather than an unidentifiable knee (§5.2).
2. A negative result on a pre-registered phoneme-class hypothesis (§5.3).
3. Evidence that accent conversion needs more lookahead than accent-faithful transcription at matched capacity (§5.4).
4. Two patches that make a masked SSL encoder actually causal, with a truncation proof (§3.2).
5. Three documented measurement failures, each of which produced plausible output before being caught (§6).

2 Related work

PHONOS [1], TVTSyn [2] and DarkStream [3] are three systems from one group, all pairing FAC or voice conversion with speaker anonymisation. Two of them report more than is sometimes credited, and it is worth being precise about what they do and do not establish. TVTSyn reports CPU as well as GPU latency (≈ 132 ms and ≈ 79 ms at a 60 ms chunk), so a CPU number for streaming conversion is not unpublished, what is unpublished is a CPU latency–quality curve. DarkStream ablates its look-ahead over $\{0, 140, 280\}$ ms and selects 140 ms on the grounds that “an extra 140ms of look-ahead yields a marginal < 1 pp absolute improvement at the expense of doubling the buffering delay”. That is a lookahead ablation, on three points, scored on anonymisation accuracy rather than conversion quality.

DarkStream’s choice corroborates our saturation point from a different task. Their marginal return collapses somewhere between 140 and 280 ms; we place saturation at ≈ 240 ms on phone error rate, measured over 16 budgets rather than three. Two systems, two objectives, the same region, which is a stronger position for this paper than the absence of prior evidence would have been. StreamVC [4] reports operation “even on a mobile platform” without naming the device, core count or thread count, which makes the claim unfalsifiable. A recent survey [5] frames the problem space. Related streaming anonymisation work [6, 7, 8] shares the streaming constraint but not the accent objective. Offline FAC [9, 10] and low-latency VC [11] bracket the design space.

Our contribution is orthogonal to all of these: we characterise an axis they parameterise but do not study. We do not claim to beat any of them on quality.

3 Method

3.1 A causal phone translator as the probe

Training a full FAC pipeline per lookahead is prohibitive. We instead isolate the component the research question is about. L2-ARCTIC [12] provides, per utterance, both the canonical phone sequence a native speaker would produce (*g2p*) and the sequence the non-native speaker actually produced (*ipa*). Accent conversion is then directly trainable as

$$\text{non-native audio} \rightarrow \text{native phone sequence},$$

a CTC head over a frozen causal encoder. This is the supervised core of the task, it costs a GPU-day rather than 200 GPU-hours, and swapping the target tensor from *g2p* to *ipa* yields a matched-capacity control that differs in *one tensor*, a tighter control than AC-versus-VC.

3.2 Making the encoder causal

Masking self-attention in wav2vec2/WavLM-base [13, 14] does *not* produce a causal encoder. Two leaks remain: the positional convolution (*pos_conv_embed*, depthwise, kernel 128, symmetric padding) admits 1.28 s of future at a 20 ms frame rate; and the feature-encoder GroupNorm

Table 1: Truncation test on wavlm-base-plus. Relative L_2 change in earlier frames after deleting later audio.

configuration	relative L_2	causal?
attention mask only	1.14×10^{-2}	no
+ causal positional conv	6.00×10^{-3}	no
+ cumulative GroupNorm	2.60×10^{-2}	no
both patches	6.14×10^{-6}	yes

normalises each channel over the entire utterance, making every output frame depend on every input frame.

We verify causality by truncation: delete audio after frame t and check whether any earlier frame moved (Table 1). Note the third row: patching GroupNorm alone makes matters *worse*. A partial causality fix is not a partial improvement.

3.3 Lookahead is quantised

Lookahead is realised as $\lceil L/f \rceil$ frames at frame rate $f = 20$ ms. A knee narrower than 20 ms is not merely unmeasured but *unrepresentable*, and sampling finer than f silently creates duplicate conditions that reweight any fit. Our dense grid is therefore 20 ms, and we assert zero collisions in t_{alg} across all conditions before fitting.

4 Experimental setup

Frozen causal WavLM-base-plus, CTC phone head, 1200 steps, batch 8, chunk 40 ms, 100-frame lookback, speaker-disjoint test split of 400 utterances. Two sweeps: a 7-lookahead $\times$ 2-target $\times$ 3-seed sweep (42 runs) and a 16-lookahead $\times$ 2-seed dense sweep, run on *both* target arms (64 runs), all on an NVIDIA T4. Compute measurements use an Apple M4 and an ARM64 Linux host, CUDA-event or *perf_counter_ns* timed, 28–50 repetitions with three warm-up iterations discarded.

Noise floor. Two sweep processes briefly ran concurrently by operator error, repeating six conditions at *identical seed and configuration*. Those repeats are the only direct measurement of fixed-seed run-to-run noise available, and a multi-seed design cannot produce it: pairing within seed cancels seed effects but not this residual. Differences ranged 1.5×10^{-4} to 6.5×10^{-3} PER, giving $\sigma = 0.0022$ and a 2σ resolvable threshold of **0.0044 PER**. Every effect below is gated on it.

5 Results

5.1 Compute is flat in lookahead

Across $L \in [0, 640]$ ms at chunk 40 ms on the M4, t_{alg} spans 40–680 ms ($17\times$) while median per-chunk compute moves from 43.23 ms to at most 44.06 ms, a range of -0.3% to $+1.9\%$. On ARM64 the span is 0 to $+7.6\%$. The two terms of Eq. 1 are therefore independently actionable, which is the practical justification for reporting them separately.




Figure 1: Dense sweep, conversion arm, 16 lookaheads × 2 seeds. **Left:** smooth curvature; the shaded band spans the breakpoint estimates produced by three defensible treatments of $L=0$, none identifiable. **Right:** marginal return per added 20 ms frame against the measured 2σ floor; red bars are below it.

5.2 RQ1: an exchange rate and a saturation point, not a knee

Conversion-arm PER falls monotonically from 0.4256 at $L=0$ to 0.1561 at $L=640$ ms. Fitting $L \geq 20$ ms on $\log_2 L$ gives

$$\Delta\text{PER} = -0.0452 \text{ per doubling}, \quad R^2 = 0.983, \quad (2)$$

over 15 points.

There is curvature, and it is not localisable. BIC prefers a two-segment fit under all three axis treatments we tried ($\Delta\text{BIC} = -22.2, -31.2, -36.8$), which by the usual thresholds [15] is strong evidence. But the breakpoint lands at 40, 180 and 180 ms respectively, and every bootstrap 90% interval spans about three octaves. Sixteen points did not rescue a knee. The reason is mechanical and worth stating: on a $\log_2(L+1)$ axis the $L=0 \rightarrow 20$ ms gap is 4.39 units while every other gap is ≈ 1.0 , so a piecewise fit places a breakpoint just past that gap whatever the data does.

What we report instead is the lookahead beyond which marginal return falls below the measured floor. Per added 20 ms frame, gains run +0.056 (0→20 ms), +0.021 (60→80), +0.0045 (180→200), and are below 0.0044 for every step past 240 ms. We therefore report **saturation at ≈ 240 ms**, read as 200–280 ms rather than a sharp value: the 140→160 and 180→200 ms steps are already only 1.0–1.5× the floor. Unlike a knee, this quantity depends only on the floor, not on an axis choice.

The cost of a 40 ms budget. PER at $L=40$ ms is 0.3275; at 240 ms it is 0.1973. A 40 ms budget therefore forgoes **39.7%** of the achievable relative reduction, rising to 52.3% against $L=640$ ms. “Current budgets are under-provisioned” is thus supportable as diminishing-returns geometry with a measurable saturation point, not as a knee.

5.3 RQ2: manner class is the wrong axis

We pre-registered, before any of this data existed, that vowels and approximants would break first (formant trajectories over 80–250 ms) while stops, fricatives, affricates and nasals would survive short lookahead (locally determined gestures). Testing this on 295,785 reference phones, with each substitution or deletion charged to the class of the *reference* phone:

Group means differ in the predicted direction (0.676 vs 0.636), but that is the weakest available reading and

Table 2: Conversion-arm PER against lookahead (mean of 2 seeds) and marginal return per added 20 ms frame, in units of the measured $2\sigma=0.0044$ floor. Below $1.0\times$ an additional frame is not measurably useful.

L (ms)	PER	gain/frame	$\times$ floor
0	0.4256	—	—
20	0.3701	0.0556	12.7
40	0.3275	0.0426	9.8
60	0.2963	0.0312	7.2
80	0.2754	0.0209	4.8
100	0.2564	0.0190	4.4
120	0.2446	0.0118	2.7
140	0.2269	0.0176	4.0
160	0.2206	0.0064	1.5
180	0.2106	0.0099	2.3
200	0.2062	0.0045	1.0
240	0.1973	0.0044	1.0
280	0.1891	0.0041	0.9
320	0.1823	0.0034	0.8
480	0.1649	0.0022	0.5
640	0.1561	0.0011	0.3

Table 3: Relative error-rate reduction, $L=0 \rightarrow 640$ ms. B = pre-registered as breaking first, S = as surviving. The predicted grouping does not order the data.

	class	rel. gain	ref. phones
S	nasal	0.716	4 020
B	approximant	0.700	4 575
S	fricative	0.687	7 659
B	vowel (mono)	0.673	13 008
B	vowel (diph)	0.654	4 824
S	stop	0.615	7 071
S	affricate	0.525	1 098

three stronger checks fail. The effect size is 0.66: the between-group difference (0.040) is smaller than the pooled SD across classes (0.060). The surviving group’s *internal* spread is 0.190, 4.7× the between-group difference. And 5 of 12 pairwise orderings are violated, with nasals outgaining all three breaks-first classes. By L1, only 3 of 6 support even the directional reading and the gain ratio straddles unity (0.92 Chinese to 1.23 Vietnamese).

H2 is not supported. The finding is not that lookahead fails to help vowels, every class improves 52–72% relative, all above the floor. The finding is that manner class does not predict which sounds need right context.

Reweighting each error by what it costs acoustically does not rescue it. PER charges every error exactly 1.0, so a natural objection is that the grouping is fine and PER is simply mis-weighting it: a substituted vowel and a substituted stop cannot sound equally wrong. We tested that directly. Each utterance’s canonical phone string is synthesised with a fixed voice and force-aligned to itself with an espeak-IPA CTC model, giving a time span per reference phone; the hypothesis is synthesised, DTW-aligned to the reference, and the spectral distortion inside each phone’s own span is charged to that phone’s class. The measure is validated before it is used: phones the model got wrong carry 1.72× the distortion of phones it got right (86.8 vs 50.5), so it does detect errors where they occur, and mean distortion falls monotonically with lookahead (71.4 to 46.9,

Table 4: H2 under acoustic weighting: relative reduction in per-phone distortion, $L=0 \rightarrow 640$ ms, 40 utterances, 1451 reference phones per condition. B = pre-registered as breaking first, S = as surviving. The sequence-metric column is computed on the same utterances, so the two differ only in weighting.

	class	acoustic	sequence	phones
B	approximant	0.386	0.825	153
B	vowel (mono)	0.379	0.712	448
S	nasal	0.376	0.784	128
S	fricative	0.372	0.745	272
B	vowel (diph)	0.321	0.667	166
S	stop	0.305	0.663	240
S	affricate	0.257	0.526	44
effect size (need ≥ 1.0)		0.756	0.607	
ordering violations		2/12	4/12	

−34%).

The verdict is unchanged (Table 4). Acoustic weighting does move in H2’s favour, effect size $0.607 \rightarrow 0.756$ and ordering violations $4/12 \rightarrow 2/12$ against the sequence metric on the same utterances, but not far enough to clear the pre-registered threshold of 1.0, and the failure has the same shape as before: the “survives” group’s internal spread (0.120) is $3.5\times$ the difference between the groups (0.034), and a nasal again out-gains a diphthong. So “PER mis-weights the classes” is no longer an available explanation for the sequence-level refutation.

One explanation does remain open, and it is not the one this test closes. Both signals here are TTS renderings of phone strings, so the synthesiser produces canonical formant trajectories for whatever string it is given. This measures the cost of getting a phone’s *identity* wrong; it says nothing about a wrong trajectory for a phone whose identity is right. That still needs real converted audio against a native rendition (§8). The sample is also far smaller than the sequence-level test, 1451 phones per condition against 295,785, and is weighted accordingly.

5.4 RQ3: conversion needs more lookahead than transcription

Identical architecture, capacity, data order, seed and step count; the two arms differ only in the target tensor. End-point gain $L=0 \rightarrow 640$ ms is -0.2695 (63.2% relative) for conversion against -0.1893 (42.8%) for accent-faithful transcription, a paired difference of $+0.0802$ (sd = 0.0035, $t = 39.7$).

The dense grid replicates this, and sharpens it. Running the full 16-lookahead $\times$ 2-seed sweep on the transcription arm as well puts both arms on the same dense grid (Table 5). The endpoint gains reproduce the 7-point result almost exactly, 63.3% for conversion against 43.5% for transcription, against 63.2% and 42.8% before, on an independent, denser sample. The exchange rates separate cleanly:

conversion -0.0452 vs transcription -0.0292 PER per doubling,

$R^2 = 0.983$ and 0.984 respectively, so **conversion buys $1.55\times$ as much per doubling of lookahead.**

Table 5: Both arms on the dense grid, mean of 2 seeds. The gap is the transcription arm’s PER minus the conversion arm’s: near zero at $L=0$ and widening monotonically with lookahead.

L (ms)	conversion	transcription	gap
0	0.4256	0.4427	+0.0171
20	0.3700	0.3896	+0.0196
40	0.3275	0.3600	+0.0325
80	0.2754	0.3271	+0.0517
160	0.2206	0.2908	+0.0702
240	0.1973	0.2755	+0.0782
320	0.1823	0.2666	+0.0843
640	0.1561	0.2501	+0.0939
exchange rate (PER/doubling)			
conversion	-0.0452	$R^2 = 0.983$	
transcription	-0.0292	$R^2 = 0.984$	

The mechanism is visible in the gap rather than the levels. At $L=0$ the two arms are nearly the same task (0.4256 vs 0.4427, a gap of 0.0171); by $L=640$ ms the gap is 0.0939, **$5.5\times$ wider**. It widens in 13 of 15 adjacent steps, and the transcription arm is worse at 16 of 16 lookaheads, every one of them by more than the 0.0044 noise floor. Lookahead, not accuracy, is what separates them, which is the claim H3 makes.

Two secondary points fall out. Both arms saturate at the same ≈ 240 ms, and neither yields a locatable knee (breakpoints move to [40, 240, 160] ms across axis treatments for transcription, against [40, 180, 180] for conversion), so §5.2’s structural conclusion is not an artefact of the conversion arm. And the transcription arm’s own curve is just as log-linear, so the shallower slope is a smaller effect of lookahead, not a noisier measurement of the same one.

The claim is carried by the *preference margin*, PER against the other arm’s labels minus PER against its own, which grows monotonically in 3/3 seeds for conversion ($+0.032$ at $L=0$ to $+0.099$ at 640 ms; 18/18 adjacent step $\times$ seed comparisons positive) while the transcription control is 0/3 monotone and drifts *down* (-0.0103). Building the claim on a difference rather than a level turned out to matter (§6).

5.5 t_{buf} : jitter measured, I/O not

On the M4, two independent 30 s duplex runs at 40 ms blocks give a required jitter buffer of 0.14 and 0.08 ms, with p99 callback lateness ≈ 0.06 ms, zero over/underruns and clock drift below $1.3 \times 10^{-4}\%$. The jitter term is **negligible**.

The I/O term is **not measured**, and we report why rather than substituting a number. An acoustic loopback failed on every attempt: the captured level was invariant to output amplitude (peak 0.00505 at amplitude 0.50 versus 0.00468 at 0.95), meaning the speaker contributes no measurable energy to the microphone, consistent with the operating system’s echo cancellation, which scales with its reference and therefore cannot be defeated by a louder probe. The driver-reported figure is not a substitute: its excess over the requested block is *identically* 201.69 ms at both 10 and 20 ms blocks and then roughly doubles, i.e. it is queueing, not con-

verter delay.

6 Three measurement failures

Each of the following produced plausible output before being caught. We report them because each is a trap the next person will meet.

A padding bug rewrote the curve’s shape. A masked block handed a zero-padded hidden-state tensor to a depth-wise convolution and feed-forward that do not respect the key-padding mask. Fixing it was *not* a constant offset: the correction grows with lookahead, from -0.014 PER at $L=0$ to -0.086 at 320 ms, a spread of $16.4\times$ the noise floor. The buggy curve therefore *understated* the benefit of lookahead, biasing the central quantity of RQ1. A constant offset would have cancelled in every within-run comparison; this did not. The preference margin (§5.4) was unaffected, because the bug moved both PERs together and the margin is a difference.

A breakpoint estimator returns a breakpoint whether or not one exists. Δ BIC cannot distinguish “there is a knee” from “piecewise fits a curve better than a line”. Identifiability must be tested by perturbing the x-axis and bootstrapping the location. Our synthetic check reproduces the artefact exactly: on a curve that is log-linear for $L \geq 20$ with $L=0$ off the extrapolation, including $L=0$ yields Δ BIC = -39 and a confident “knee at 40 ms”; dropping it yields $+5$ and no knee.

An analysis that buckets unmatched symbols into “other” will produce confident output on the wrong alphabet. Our phoneme-class map was keyed on ARPAbet; the sweep emits IPA. Applying one to the other sent 57% of reference tokens, including *every vowel* and $/r/$, to “other”, and still printed per-class slopes, an H2 verdict and a by-L1 table. Assert coverage; do not infer it from the absence of an error.

7 Limitations

Every number here is phone error rate, not perceived quality. No perceptual evaluation was run, so nothing supports a claim about naturalness or audible accentedness: the saturation point says marginal *PER* gain falls below the noise floor, not that listeners stop noticing improvement. This is the limitation a reader should weigh most heavily.

We report why, because “unfunded” would be the wrong reason. Stimuli can be built from the per-condition outputs and were, but they cannot carry the judgement a listening test needs, for two independent reasons. The phone sequences carry no word or stress structure, so even the ground-truth ceiling synthesises as one unsegmented stream, $/\text{a}\text{n}\text{e}\text{v}\text{r}\ \text{a}\text{l}\text{a}\text{u}\text{o}\text{w}\text{a}\text{l}\text{k}\text{æ}\text{n}\text{t}\text{b}\text{i}\dots/$ for “*I never allow what can’t be changed...*”. And a system at 0.16–0.42 PER does not emit recognisable words to compare. An informal session confirmed both: the listener could not identify what was being said, and scored at chance on ground-truth anchors. **A funded rater study on these stimuli would not have measured accentedness at any budget.** What perceptual evaluation needs here is converted audio from a synthesis pipeline, which is future work, not a purchase.

8 Remaining work

Four items are outstanding, listed with what each would settle.

(i) **H2 on real converted audio.** The identity-error half of this is now done (§5.3): forced alignment charges each error its measured acoustic cost, and the grouping still fails. What remains is the trajectory half, which a TTS cannot answer because it renders canonical trajectories by construction, it needs a trained synthesis pipeline whose output can be aligned against a native rendition. Note also that a VITS-class synthesiser is globally coupled, so whole-utterance distortion cannot be charged to one phone; the per-phone span from forced alignment is what makes the measurement local at all (`docs/H2_AUDIO_DOMAIN_BLOCKED.md`).

(ii) t_{buf} **I/O on a platform without echo cancellation.** A wired loopback, or measurement on the ARM64/Raspberry-Pi class this paper’s deployment argument is about, would close the one term of Eq. 1 that remains unmeasured.

(iii) **The hardware sweep beyond two architectures.** Compute is currently characterised on Apple Silicon and ARM64 Linux. The embedded class is projected, not measured.

(iv) **Perceptual evaluation.** Stimulus generation is implemented and the inputs exist; only the rater study is unfunded. Every claim in this draft is therefore scoped to PER, and the paper should not be submitted without either running some form of perceptual check or stating this limitation as prominently as §8 does.

9 Data and code availability

Everything needed to reproduce every number is public. The harness, the causality proof, the sweep configurations, the analysis scripts and their self-tests, and the raw per-condition outputs are at <https://github.com/Dipeshtripathi13/streaming-fac-lookahead>. This includes the per-condition hypothesis dumps for both target arms, which are the input to the phoneme-class and audio-domain analyses and cost a full sweep to regenerate, and the host-identity JSON recorded alongside every timing.

Speech data is L2-ARCTIC [12], obtained through the KoelLabs/L2Arctic release (3599 phoneme-annotated utterances, 24 speakers, 6 L1s). That release is CC-BY-NC-4.0. The licence is non-commercial, which is why no trained model weights are distributed; the training code and configurations are released instead, and reproduce the reported curves on a single mid-range GPU in a few hours.

10 Acknowledgements

This work used no external funding and no institutional compute. All GPU results were obtained on free-tier Google Colab (Tesla T4); all CPU results on consumer hardware. The author thanks the maintainers of L2-

ARCTIC, WavLM, sherpa-onnx and Piper, whose public releases made a single-author study of this scope possible.

[TODO: swap in the official Interspeech style file when the 2027 template is published, and re-check the page limit: this draft is 4 pages of content plus references, which meets the project's 4+2 target but not a 4+1 rule]

11 Conclusion

Lookahead buys accent-conversion quality at a stable exchange rate of ≈ 0.045 PER per doubling, with no knee to locate and measurable returns out to ≈ 240 ms, six times the budget the current state of the art uses, which forgoes about 40% of the achievable reduction. Per-chunk compute is flat across that range, so the trade-off is a modelling decision rather than a hardware one. A pre-registered phoneme-class hypothesis was refuted. We release the harness, the per-condition hypothesis dumps, and the three measurement failures that had to be caught along the way.

References

- [1] Waris Quamer, Mu-Ruei Tseng, Ghady Nasrallah, and Ricardo Gutiérrez-Osuna. PHONOS: PHOnetic neutralization for online streaming applications. *arXiv preprint arXiv:2603.27001*, March 2026.
- [2] Waris Quamer, Mu-Ruei Tseng, Ghady Nasrallah, and Ricardo Gutiérrez-Osuna. TVTSyn: Content-synchronous time-varying timbre for streaming voice conversion and anonymization. *arXiv preprint arXiv:2602.09389*, February 2026.
- [3] Waris Quamer and Ricardo Gutiérrez-Osuna. DarkStream: Real-time speech anonymization with low latency. *arXiv preprint arXiv:2509.04667*, September 2025.
- [4] Yang Yang, Yury Kartynnik, Yunpeng Li, Jiuqiang Tang, Xing Li, George Sung, and Matthias Grundmann. StreamVC: Real-time low-latency voice conversion. In *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2024.
- [5] Yurii Halychanskyi, Jianfeng Steven Guo, and Volodymyr Kindratenko. Accent conversion: A problem-driven survey of sociolinguistic and technical constraints. *arXiv preprint arXiv:2604.27281*, April 2026.
- [6] Nikita Kuzmin, Songting Liu, Kong Aik Lee, and Eng Siong Chng. Stream-voice-anon: Enhancing utility of real-time speaker anonymization via neural audio codec and language models. In *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2026.
- [7] Nikita Kuzmin, Kong Aik Lee, and Eng Siong Chng. StreamVoiceAnon+: Emotion-preserving streaming speaker anonymization via frame-level acoustic distillation. *arXiv preprint arXiv:2603.06079*, March 2026.
- [8] Waris Quamer and Ricardo Gutiérrez-Osuna. End-to-end streaming model for low-latency speech anonymization. *arXiv preprint arXiv:2406.09277*, June 2024.
- [9] Guanlong Zhao, Shaojin Ding, and Ricardo Gutiérrez-Osuna. Converting foreign accent speech without a reference. *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, 29:2367–2381, 2021.
- [10] Shaojin Ding, Guanlong Zhao, and Ricardo Gutiérrez-Osuna. Accentron: Foreign accent conversion to arbitrary non-native speakers using zero-shot learning. *Computer Speech & Language*, 72:101302, 2022.
- [11] Konstantine Sadow, Matthew Hutter, and Asara Near. Low-latency real-time voice conversion on CPU. *arXiv preprint arXiv:2311.00873*, November 2023.
- [12] Guanlong Zhao, Sinem Sonsaat, Alif Silpachai, Ivana Lucic, Evgeny Chukharev-Hudilainen, John Levis, and Ricardo Gutiérrez-Osuna. L2-ARCTIC: A non-native English speech corpus. In *Interspeech*, pages 2783–2787, 2018.
- [13] Sanyuan Chen, Chengyi Wang, Zhengyang Chen, Yu Wu, Shujie Liu, Zhuo Chen, Jinyu Li, Naoyuki Kanda, Takuya Yoshioka, Xiong Xiao, Jian Wu, Long Zhou, Shuo Ren, Yanmin Qian, Yao Qian, Michael Zeng, Xiangzhan Yu, and Furu Wei. WavLM: Large-scale self-supervised pre-training for full stack speech processing. *IEEE Journal of Selected Topics in Signal Processing*, 2022.
- [14] Alexei Baevski, Henry Zhou, Abdelrahman Mohamed, and Michael Auli. wav2vec 2.0: A framework for self-supervised learning of speech representations. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.
- [15] Robert E. Kass and Adrian E. Raftery. Bayes factors. *Journal of the American Statistical Association*, 90(430):773–795, 1995.